KiraOrder for Orbyte — Product Requirements Document

Document control

Field	Value
Product	KiraOrder for Orbyte
Team	Falcons — Xiamen University Malaysia
Track	Domain 2 — Orbyte Smart SME Challenge (SDG 8.2, 8.3)
Document version	v2.0 (Final round)
Status	Final round v2.0 — submission copy
Prepared	2026-10-31; refreshed 2026-11-01, UTC+8
Evidence baseline	Repository main @ 809832b (2026-11-01); requirements read at f347ed6 (2026-10-31 20:18:10 +0800, 257 commits). After the earlier reads, commits add frontend 1.5.2, the F03/F04 front-end work and the review-display merge d4915e7 (100/100 tests); Pages serves 1985f0bc from f347ed6, so d4915e7 is not deployed. Earlier reads: 7bf4b47 (11:01), 53996ff (11:52), a0628d3 (12:04), 8365c928 (12:22). The a0628d3 commit fixes the unit-blocking gap described in §7.1 and §9.
Supersedes	Phase 1 preliminary proposal, 2026-10-25
Structure note	**The organisers provide no PRD template.** This structure is a team convention chosen for reviewability. It is not an official SDC form and does not claim to be one.

How to read the status tags. This document separates what the product is *designed* to do from what has actually been *observed*. Every requirement carries one tag.
[VERIFIED] observed in a recorded acceptance run, with an owner or reviewer sign-off.
[BUILT] implemented, merged and covered by automated tests, but not accepted through a recorded business scenario.
[MERGED, NOT DEPLOYED] code is on main but the running deployment still serves the previous build.
[DESIGNED] specified here; implementation or acceptance is incomplete.
[NOT DONE] committed to in Phase 1 and still outstanding. Stated as a gap, not softened.

1 Purpose and scope

This document specifies the final product: one recurring operational chore for one Malaysian wholesale SME, delivered end to end, with the human decision points made explicit.

It is the requirements counterpart to the Final Project Report. It states what the system must do, what it must refuse to do, and — separately — what has been demonstrated. It does not restate the pitch.

Section 9 records every material change from the Phase 1 proposal and why it was made. Section 10 states the acceptance conditions together with the known limitations that a judge should test against.

2 Problem and target users

2.1 The user we build for

Named user: the order clerk at **Sinaran Food Supply**, a micro food wholesaler in the Klang Valley. One clerk takes the evening orders; one owner approves anything that costs margin.

Recurring chore: a mixed-language WhatsApp order arrives at 18:04. Before anything can be picked, the clerk must decide what the message does *not* say.

A synthetic example, written by us:

Formula

```
2 bag Gula, 20kg bwg bsr India,
5 包美禄粉, kalau ada…
"kalau RM2.80 boleh?"
```

Six facts are needed and none of them are in the message: which Gula pack size (1 kg or 2 kg), which tier price this customer buys on, which egg grade, whether the onion is in stock and how much, what `kalau ada` means here (reduce the line, or drop it), and whether RM2.80 sits below her tier price.

2.2 What the chore costs today

Phase 1 recorded two figures and labelled them honestly: **≈210 seconds median handling time per five-line order** and **up to 30 % transcription or stockout discrepancy**. Both were **team planning estimates, not measurements**. They are retained here for continuity and are flagged as unvalidated.

`[NOT DONE]` No measured replacement for either figure is available. The baseline package exists — six participants, five cases each, thirty planned runs — and all thirty raw rows are `not_started` with no elapsed time recorded; the first-use sheet and the session log are empty. That is a statement about the package as it stands, not proof that no measurement was ever attempted elsewhere. See §10.4. No time-saving or error-reduction claim is made in this product's materials while no measured result is available.

2.3 Why existing answers are insufficient

Existing approach	What it does	Why it stops here
Copy the message into a spreadsheet	Faithful transcription	Nothing is verified; the pack size, tier price and stock are still unknown
A translating bot	Produces a clean structured list	The list looks perfect and can still be wrong. `kalau ada` and pack size are not language problems

Existing approach	What it does	Why it stops here
A flag-everything checker	Marks suspicious lines red	A red flag is a question, not an answer; alert fatigue follows

For a shop selling one product at one price, any of the three is enough. Sinaran's catalogue has four pack sizes, six tier prices and a stock board on the wall.

2.4 The design consequence

The product does not try to be certain. It tries to be **structurally honest**: it settles what the data settles, and where the data does not settle a question, it stops and names the person who can.

3 Product principles

1. **AI proposes; rules and people decide.** The model extracts candidate products, quantities, units, packs and conditional phrases. It cannot set a price, approve an order, or generate a delivery order.
2. **Never guess silently.** Where a field cannot be evidenced, the line becomes `UNKNOWN` with an explicit reason code, not a plausible value.
3. **Stop at the right person.** A missing pack size goes back to the clerk; a below-tier price goes to the owner; a stock problem goes to the clerk with the count. `[VERIFIED]`
4. **The approval is enforced, not drawn.** A blocking issue prevents approval at the API, not merely by disabling a button. `[VERIFIED]`
5. **Every consequential action is attributable.** Save, approve, reject, deduct and generate each write an actor and a timestamp. `[VERIFIED]`

4 Core business flow

The six-step journey required by the Domain 2 final guide, mapped to what is built:

Step	Requirement	Implementation	Status
1 Trigger	A business request starts the work	The clerk pastes the WhatsApp message into **New order** and selects the customer	`[VERIFIED]` manual trigger. Automatic triggers are not built
2 Connected data	Relevant business information is pulled together	Tenant catalogue, customer aliases, tier prices, pack-size candidates, inventory, and customer phrase memory	`[VERIFIED]`
3 Decision logic	Rules or AI decide what happens; exceptions go to a person	Exact alias resolution and rules settle price, stock and pack size; the language model produces candidates and spans only	`[VERIFIED]` including explicit exception routing

Step	Requirement	Implementation	Status
4 Human approval	A person with the right role approves the consequential action	Routine orders: CLERK approves. Below-tier price: OWNER decides. Rejection stops the order	`[VERIFIED]`
5 Visible outcome	The action completes and the activity log shows what happened	Delivery order generated, inventory deducted, activity and per-field change record available	`[VERIFIED]`
6 Second business	How a different SME adopts the same product	A second tenant, hardware and cleaning supplies, runs on the same code with different seed configuration	`[PARTIAL]` See §8

5 Functional scope

5.1 In scope

ID	Function	Detail	Status
F-01	Mixed-language order capture	Malay, English and Chinese in one message; customer-specific product names and aliases	`[VERIFIED]`
F-02	Candidate generation with spans	Model returns candidate product family, quantity, unit, pack size and conditional span, restricted to supplied family IDs or `UNKNOWN`	`[VERIFIED]` on 23 development cases and 6 contract samples
F-03	Pack-size disambiguation	Multiple valid pack sizes are presented unselected; approval is blocked until one is chosen **and saved**	`[VERIFIED]`
F-04	Price and stock computation	Backend rules compute tier price, line totals and stock availability; the model never sets a price	`[VERIFIED]`
F-05	Blocking issues	Missing unit, missing pack size, insufficient stock and unresolved ambiguity each produce a named blocking issue	`[VERIFIED]`
F-06	Price exception escalation	A below-tier price creates an approval request; the draft cannot be edited while awaiting the owner	`[VERIFIED]`
F-07	Owner approve or reject	Owner sees a saved snapshot, approves, or rejects with a mandatory reason	`[VERIFIED]`

ID	Function	Detail	Status
F-08	Delivery order generation	Approval generates a read-only delivery order bound to the saved draft version	`[VERIFIED]`
F-09	Inventory transaction	Approval deducts stock in the same transaction; re-approval cannot deduct twice	`[VERIFIED]`
F-10	Activity log and field change record	Who, what, when, decision and outcome; per-field before and after	`[VERIFIED]`
F-11	Restricted warehouse view	Warehouse finds a delivery order by number and sees 4 top-level fields only; no prices, no contact details, no original message	`[VERIFIED]`
F-12	Manual recovery	A clerk can add a line, reselect a SKU, or change a price target, with confirmed SKU, integer quantity and `MANUAL_SELECTED` provenance	`[BUILT]` — backend reviewed to PASS at the API scope; front-end F03 implemented (78/78 frontend tests, browser acceptance on synthetic data) but **not yet independently reviewed**
F-13	Read-only product catalogue	`GET /api/v1/catalog/products`, at most 100 active products per tenant, explicit `truncated` flag	`[BUILT]`
F-14	Bilingual customer replies and clarification copies	Reply drafts are copied to the clipboard; the system does not send WhatsApp messages	`[VERIFIED]`
F-15	Unsaved-change and version protection	Unsaved edits block approval; version conflicts and stock changes are detected at decision time	`[VERIFIED]`
F-16	Isolated acceptance environment	Two independent PostgreSQL clusters allow write-path acceptance without touching the production demo database, then reset	`[VERIFIED]`

5.2 Explicitly out of scope

Not built	Reason
Sending or receiving WhatsApp messages	The product copies drafts; it does not connect to WhatsApp. Messaging integration would require business verification we do not have

Not built	Reason
Real payments, statutory filings, legal or financial advice	Prohibited by the challenge rules
Any real customer, supplier or company record	Only team-created synthetic data is used
Automatic triggers (new order detection, scheduled runs)	Not built. The trigger is a manual paste. This is stated as a gap, not implied as a feature
A complete private TEST32 workflow suite	12 workflow cases remain `NOT_RUN`
Multi-currency, multi-warehouse, credit control	Roadmap only

6 Roles, approval and activity log

6.1 Role matrix

Identity, tenant and role come only from the bearer token's user record. Request bodies, query parameters and headers cannot override them. `[VERIFIED]` by a 41-case API matrix.

Capability	CLERK	OWNER	WAREHOUSE
Create and parse a draft	yes	—	—
Edit lines and resolve issues	yes	—	—
Approve a routine order	yes	yes	—
Submit a price exception to the owner	yes	—	—
Approve or reject an exception	—	yes	—
Read the product catalogue	yes	yes	denied (403)
Read a delivery order, full view	yes	yes	—
Read a delivery order, picking view	—	—	yes, 4 fields only
Read the activity log or per-field audit	yes	yes	denied (403)

Notes that must not be lost:

- Role switching inside **Local Demo** is a browser-only convenience. In connected API mode, permissions come from the authenticated backend session; the demo role selector grants nothing. `[VERIFIED]`
- The warehouse view is an explicit field projection, not a hidden-UI illusion: the same delivery order returns 14 top-level fields to CLERK and 4 to WAREHOUSE, and the total amount does not appear anywhere in the warehouse response body. `[VERIFIED]`

6.2 Approval enforcement

Requirement	Behaviour	Status
A blocking issue prevents approval at the API	Approving a draft with unresolved issues returns `400 DISAMBIGUATION_UNRESOLVED` and generates no delivery order	`[VERIFIED]`
Insufficient stock prevents approval	`400 INSUFFICIENT_STOCK`; the draft stays `READY` after a failed attempt	`[VERIFIED]`
Stock changed after submission is detected at decision time	`409 STOCK_CHANGED`; the approval request stays `PENDING`; no delivery order; no negative inventory	`[VERIFIED]`
Double approval is impossible	Second approval returns `409 ALREADY_APPROVED`; with four concurrent approvals exactly one succeeds and three receive 409, one delivery order is created, stock moves once	`[VERIFIED]`
A rejected order cannot be acted on again	Further decisions return `409 DRAFT_REJECTED`; no delivery order; stock unchanged	`[VERIFIED]`
The last unit cannot be sold twice	One approval succeeds, one returns `400 INSUFFICIENT_STOCK`, stock 1 → 0, one delivery order	`[VERIFIED]`

6.3 Activity log

Every entry is written when the event happens, not reconstructed.

`[VERIFIED]` Example from a recorded run: a save recorded at 21:52 and an inventory deduction plus approval at 21:55, all attributed to the same synthetic clerk. A rejection run recorded the draft version at 01:45, the owner-approval request by the clerk at 01:46, and the owner's rejection with its full reason at 01:53.

Available surfaces: the order activity timeline, the per-field change record, and the delivery-order audit detail.

7 Exception handling and manual recovery

7.1 Exception routing

Exception	Detection	Routed to	Resolution	Status
Pack size not stated	Two or more valid candidates, none selected	Clerk	Select one, save, then approve	`[VERIFIED]`
Product not identified	Row returned as `UNKNOWN` with `LOW_CONFIDENCE_MATCH`	Clerk	Choose from candidates with recorded provenance, or ask the customer	`[VERIFIED]` parse path

Exception	Detection	Routed to	Resolution	Status
Unit not evidenced in a multi-unit family	`uom = null` with `FIELD_NOT_EVIDENCED`; `sku_id` left null and a `MISSING_UOM` blocking issue raised	Clerk	Choose the unit explicitly, then save	`[VERIFIED]` parse path at `555f0fa`; save path hardened 2026-10-31
Conditional phrase span truncated	Extracted span shorter than the source clause; **no issue is raised**	—	Nothing in the current build flags it; the clerk must read the line	`[KNOWN GAP]` 4 of 8 spans in the evaluated sample
Insufficient stock	Counted against inventory at save and at approval	Clerk	Reduce the quantity, or wait for stock	`[VERIFIED]`
Stock changed while awaiting the owner	Inventory recheck at decision time	Owner	Re-open the draft; the approval must be re-derived	`[VERIFIED]`
Below-tier price	Computed against the customer tier	Owner	Approve, or reject with a reason	`[VERIFIED]`
Ambiguous conditional phrase	Conditional span extracted and shown	Clerk, then the customer	Clerk resolves from the count or asks the customer	`[VERIFIED]`
Duplicate or repeat approval	Idempotency key and state check	System	409, no second delivery order	`[VERIFIED]`

7.2 Manual recovery

The recovery path exists because the honest answer to "what if the extraction is wrong" must be a workflow, not an apology.

Requirements: `[BUILT]` — backend implemented, merged and deployed, and reviewed to PASS at the API scope by a team member who did not write it. The front-end picker and add-row surface is a separate task and is **not** signed off.

What that review does and does not cover, because the difference matters if a judge quotes it: the reviewer confirmed the pull-request head and merge, both source hashes, and CI, and confirmed that manual recovery, price-condition rebinding and the below-tier owner requirement hold. The reviewer **did not** independently re-run the 151 locally reported PostgreSQL cases, **did not** independently verify the running deployment, and the front-end page was **not** in the reviewed change set. The review was received after the merge and is recorded as such, not backdated. One non-blocking defect is retained and known: invalidated metadata may keep a `price_basis` value and a related delivery-order number.

- A new line requires a selected SKU, a confirmed integer quantity, and an explicit `MANUAL_SELECTED` provenance.
- A manual selection may not present itself as a model result. The match source changes from the fallback value to `MANUAL_SELECTED`, and that change is visible in the change record.
- Adding a line requires saving first so the server issues a row identifier, then binding price conditions to the new version.
- Deleting a line or changing a price target preserves the original condition, invalidates the affected confirmations and row references, and requires reconfirmation.
- A below-tier manual price still requires owner approval.
- No new endpoint, migration, dependency, seed or frozen contract change was introduced; recovery reuses the existing save and approve routes.

`[DESIGNED]` Front-end: product picker, add-row entry point, unsaved-edit protection and browser-level acceptance. Do not claim this as delivered.

8 Orbyte fit and second-SME reuse

8.1 Fit

Orbyte is a shared business hub that brings functions, information and services together. KiraOrder is deliberately scoped as a **connectable order-intake capability**, not a standalone tool: it consumes a tenant catalogue, customer records and inventory as configuration, and produces delivery orders and audit records that other hub services could read.

It does **not** reproduce Orbyte's existing HR or administrative functions. The capability it adds is order-intake judgement: converting an ambiguous message into a verified, approved order.

8.2 Reuse across SME types

Second tenant: **Kilat Hardware & Cleaning Supplies** — hardware and cleaning supplies, a different product category from food wholesale.

Configuration	Tenant 1 — Sinaran Food Supply	Tenant 2 — Kilat Hardware & Cleaning Supplies
Products	40	16
Inventory units	795	450
Users	3	5
Customers	8	5
Product family aliases	50	30

Configuration	Tenant 1 — Sinaran Food Supply	Tenant 2 — Kilat Hardware & Cleaning Supplies
Unit aliases	26	30
Seed version	`KB-T1-20261022-r1`	`KB-T2-20261022-r1`

What a second SME changes: seed data, catalogue, aliases and tenant policy settings. What it does not change: the application code, the API contract, the approval model or the audit schema. Tenant isolation is enforced at the API: a Tenant 2 token requesting a Tenant 1 delivery order receives 404, and neither `?tenant_id=` nor an `X-Tenant-Id` header changes the result. `[VERIFIED]`

`[DESIGNED]` A complete second-tenant business walkthrough — parse, approve, deliver for Kilat — has **not** been recorded and signed off. What is verified is the data exists and the isolation holds. Do not describe Tenant 2 as a fully demonstrated second deployment.

8.3 Demonstration paths — decision open

The team has not yet fixed which path carries the live demonstration. This document does not decide it. The three candidate paths carry different evidence, and mixing them is the most likely way to over-claim on stage.

Path	What it can show	What it cannot establish	What it depends on
Connected back end (API mode)	Persistence in PostgreSQL; identity and tenant resolution from the session; the approval gate at the API; the six concurrency scenarios; the 41-check permission matrix; field projection; the activity log written by the server	Model quality, unless the worker is warm and the run is recorded — a fallback parse is not a model result	A role-bound session token handed over privately; a reachable deployment; the demo database in a known state
Browser Local Demo	Interaction inside one browser: draft review, blocking checks, approval and rejection flow, change record, warehouse view layout	Server-side approval enforcement, database persistence, API field cropping, tenant isolation, model behaviour — none of these are exercised	Nothing but a browser. Whether it runs with the network disconnected has **not** been tested
Backup recording	A repeatable narrative of whichever path was recorded, with a known build identifier	Anything not on screen in that recording; and it cannot be described as live	A recorded file and a link. **No recording exists yet**

Any statement of the form "we have verified X" must name which of these three paths produced it. A demonstration that merely opens API mode does not by itself verify API-scope behaviour; that requires a recorded run.

9 Changes since the Phase 1 proposal

Each change is stated with the reason and the status. Changes marked "no improvement" are recorded because a judge will compare the two documents.

#	Phase 1 position	Final position	Why it changed	Status
1	Accuracy `NOT_COMPUTED`; no answer key existed	Independent answer key built; strict six-field case pass measured at **8/23 (34.8 %)**	A zero-fallback call run is not a quality result and would not survive a judge's question	`[VERIFIED]`
2	Unit handling assumed sound	**7 known unit gaps** identified (4 bulk-kg, 1 botol, 2 pail); each becomes a blocking `MISSING_UOM` rather than a guess	The measurement found them; hiding them would be a silent mismatch of our own	`[VERIFIED]`
3	The configured AI extraction path was presented as the main integration; a browser-only Local Demo existed alongside it	Same two paths. Phase 2 added cold-start evidence (2026-10-29: 180-second read timeout, five consecutive fallbacks before the worker warmed), which constrains *when* the model path can be shown live. **Which path carries the live demonstration is a team decision that is still open** — this document does not make it	A live model that can fall back for the first minutes needs a stated recovery, not an assumption that it will be warm	`[DESIGNED]` decision open
4	Human approval was already specified as a stage, with no enforcement evidence	Enforcement verified at the API against six recorded concurrency and idempotency scenarios, per scenario rather than in aggregate	Phase 1 had the design; Phase 2 supplied evidence, which is a different claim	`[VERIFIED]`
5	A restricted warehouse view was asserted; the C01 evidence boundary was limited to "no page entry" plus `permissions=[]`	41 checks across seven categories on the production API, a field-projection diff (14 top-level fields for the clerk, 4 for the warehouse) and a response field scan	The earlier position was an evidence boundary from C01, not a Phase 1 claim; UI absence is not a permission control	`[VERIFIED]`
6	Manual correction was unspecified	Manual recovery with confirmed SKU, integer quantity and explicit provenance, plus condition revalidation	"What if it is wrong" is the first question a judge asks	`[BUILT]`

#	Phase 1 position	Final position	Why it changed	Status
7	No catalogue surface	Read-only catalogue endpoint and page, tenant-scoped, 100-item cap with explicit truncation	Manual recovery and manual selection need something to select from	`[BUILT]`
8	Acceptance could only run against the shared demo database	Isolated two-cluster environment with a verified reset	Write-path acceptance without endangering the demo data	`[VERIFIED]`
8a	A missing unit blocked at parse, and the save path was assumed to behave the same	**The save path could fill the unit from the selected SKU and clear the block.** Auto-fill removed on 2026-10-31; a missing unit now blocks save and approve, with regression tests	An assumption about a second code path was tested and failed; the failure is recorded rather than quietly fixed	`[VERIFIED]` fix
9	Handling time promised as a Phase 2 measurement	**No measured value available**	No verifiable result was found in the baseline package (30 of 30 planned runs `not_started`) or supplied to this review	`[NOT DONE]`
10	First-use completion ≥80 % promised	**No measured value available**	The first-use sheet in the baseline package is empty; no verifiable result was supplied	`[NOT DONE]`
11	12 TEST32 workflow cases `NOT_RUN`	**Still 12 `NOT_RUN`**	No new workflow batch is recorded in the repository	`[NOT DONE]`
12	Independent UAT planned for Phase 2	**Not performed**	—	`[NOT DONE]`
13	Candidate confirmation and customer memory rested on indirect FALLBACK-mode evidence	**C11 signed off**: candidate confirmation in model mode plus customer-level memory write and same-phrase reuse, 7 PASS / 1 not-a-defect / 2 limitation	"What happens when the model is wrong" needed a positive answer, not only a blocking one	`[VERIFIED]`
14	Reproducibility was an untested assumption	**C13 tested and did not pass**; recorded as a known limitation. AI-unavailable handling (C10) is partial	The rubric expects the same input to give the same correct result; it does not, so we say so	`[KNOWN LIMITATION]`

Feedback basis: no written record of mentor or judge feedback was found in the repository or supplied to this review, so this document does not attribute any change to feedback. The priorities

above were set by the team from the recorded evidence — that the demonstration path must survive a cold start, that the approval must be real, and that recovery must be usable.

10 Acceptance conditions and known limitations

10.1 Acceptance conditions

A build is acceptable for the final demonstration when all of the following hold on the exact build that will be shown:

6. A routine order completes: paste → draft → review → approve → delivery order → inventory deducted → activity log populated.
7. An ambiguous pack size blocks approval until a candidate is selected **and saved**.
8. A below-tier price cannot be approved by the clerk; the owner sees it, and a rejection with a reason closes the order with no delivery order.
9. Insufficient stock blocks approval and produces no delivery order.
10. A repeated approval returns 409 and does not create a second delivery order or a second stock movement.
11. The warehouse view of the same delivery order contains no price, total or customer contact detail.
12. Every step above appears in the activity log with an actor and a timestamp.

All seven were met on recorded builds; see the evidence index.

10.2 Known limitations

ID	Limitation	Consequence
L-01	Strict six-field accuracy is 8/23 on development cases	The product is a candidate generator, not an autopilot. Any material that implies otherwise is wrong
L-02	7 unit gaps remain	Units must be confirmed by a human for bulk, botol and pail lines
L-03	**Reproducibility (C13) was tested and did not pass**: the same input in the same context did not reliably return the same result. Model calls are also single hot-state runs	No repeatability or stability claim is made anywhere in our materials
L-04	Cold start can fall back for the first several orders	The live model is not the primary demo path
L-05	TEST32 incomplete: 12 workflow cases `NOT_RUN`	No end-to-end suite claim

ID	Limitation	Consequence
L-06	No independent UAT and no real SME trial	No usability or adoption claim
L-07	Handling time, error rate and first-use completion are unmeasured	**No productivity number appears in any material**
L-08	Front-end manual-recovery surface (F03) is implemented but not independently reviewed	Do not present reselection as signed off. The backend path is reviewed at the API scope and the F03 implementation has its own browser acceptance; neither is a formal TEST32 result
L-09	Front-end review-display changes are merged but not deployed	The live site serves 1985f0bc from f347ed6 without the merged missing-tier-price diagnostics
L-10	Second tenant business walkthrough not recorded	Second-SME claims must be scoped to configuration and isolation
L-11	Deployment identity is partly inferred from publish records and diffs, not a byte-level readback of the running server	State it as recorded, not as proven
L-12	Two older pending owner requests may still be visible in the demo workspace	Clean up or label before the pitch
L-13	**Handling when the model is unavailable (C10) is only partially verified**	A degraded or unavailable model may not fail as gracefully as intended; the recorded evidence covers part of the path only

10.3 Privacy posture

Only team-created synthetic data is used. Free-text fields are masked before any model call. Session tokens live in tab memory and are never written to the repository or to screenshots. The production demo may contain synthetic customer names and synthetic phone numbers; these are not real people. `[DESIGNED]` A full PDPA retention and consent description belongs in the Technical & Tools Declaration, §6.1.

10.4 The one number that matters most

The Phase 1 proposal promised a measured comparison against a manual baseline. No measured result is available to this document: the package is prepared and empty, and no timing was supplied to this review. Until a result exists, this document makes no productivity claim and the Final Report carries an explicit placeholder.

Two things are worth separating, because they are often merged. *Not having a result* is a gap in our evidence. *Saying the team never tried* would be a claim we cannot support — the package exists precisely because the measurement was planned in detail. The honest position is: prepared, not executed as far as any record shows, and not available for this document.

11 References

- Repository: <https://github.com/Hozier123/sdc-2026-team-project> at `809832b` (requirements read at `f347ed6`; review-display merge `d4915e7`)
- Evidence index: `evidence/EVIDENCE_INDEX.md` (dated 2026-10-24; **superseded for Phase 2 results** by `evidence/phase2/2026-10-30/` and `evidence/phase2/2026-10-31/`)
- Phase 2 AI results: `evidence/phase2/2026-10-30/ai-dev-results.md`, `ai-review-summary.md`, `ai-degradation-recovery.md`
- Business acceptance: `evidence/e2e/c01/2026-10-29/`, `evidence/phase2/2026-10-30/c02/`, `evidence/phase2/2026-10-31/c05/`, `evidence/e2e/c06/2026-10-31/`
- Concurrency and permissions: `evidence/e2e/c07/2026-10-31/`, `evidence/e2e/c08/2026-10-30/`
- Manual recovery: `evidence/b03/2026-10-31/`, `docs/ops/b03-manual-recovery.md`
- Catalogue: `docs/ops/b02-catalog.md`
- Official rubric: SDC 2026 Domain 2 Finalist Guide, pp. 6–12